

# Schema Registry integration with YangKit

INSA de Lyon

Alex Huang Feng ([alex.huang-feng@insa-lyon.fr](mailto:alex.huang-feng@insa-lyon.fr))

Vivek Boudia ([vivekananda.boudia@insa-lyon.fr](mailto:vivekananda.boudia@insa-lyon.fr))

Pierre Francois ([pierre.francois@insa-lyon.fr](mailto:pierre.francois@insa-lyon.fr))

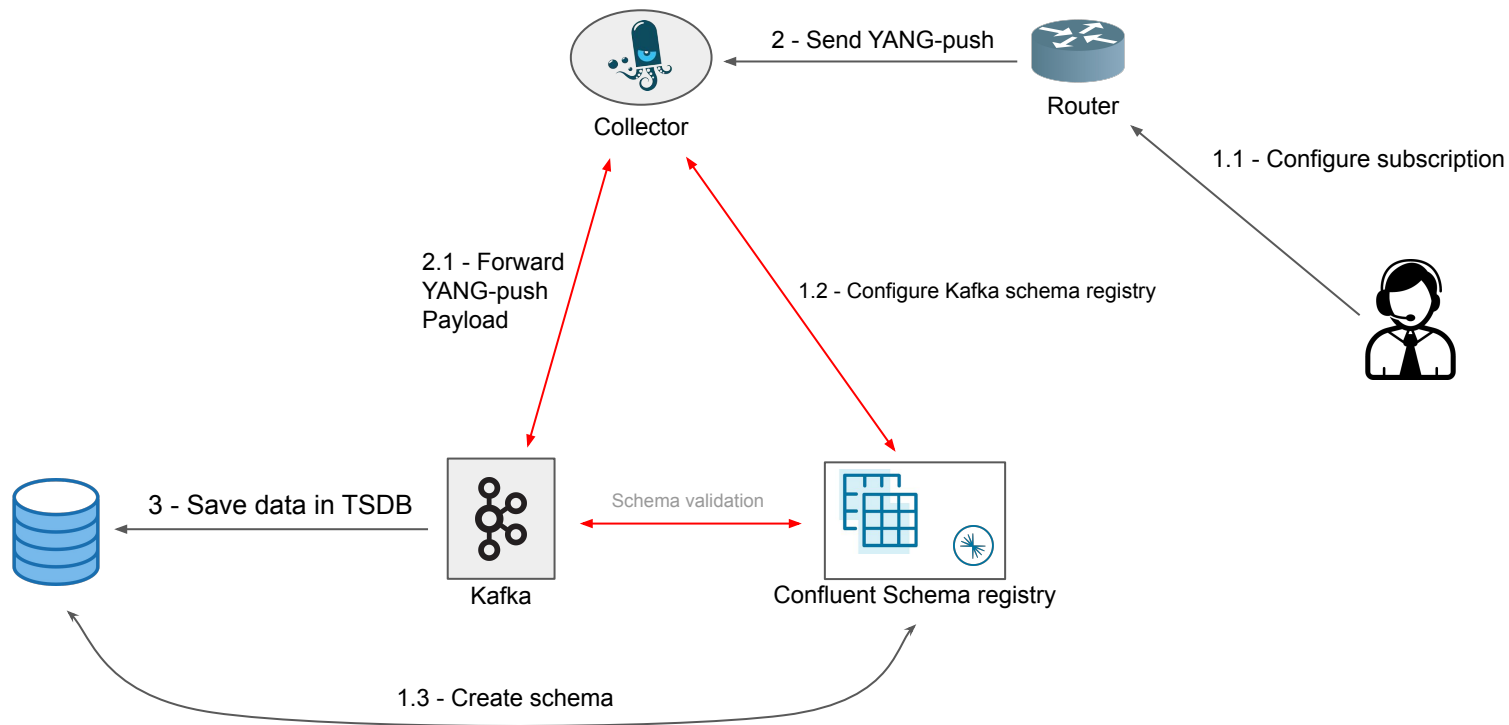
*Last Updated: 27 october 2023*

# Index

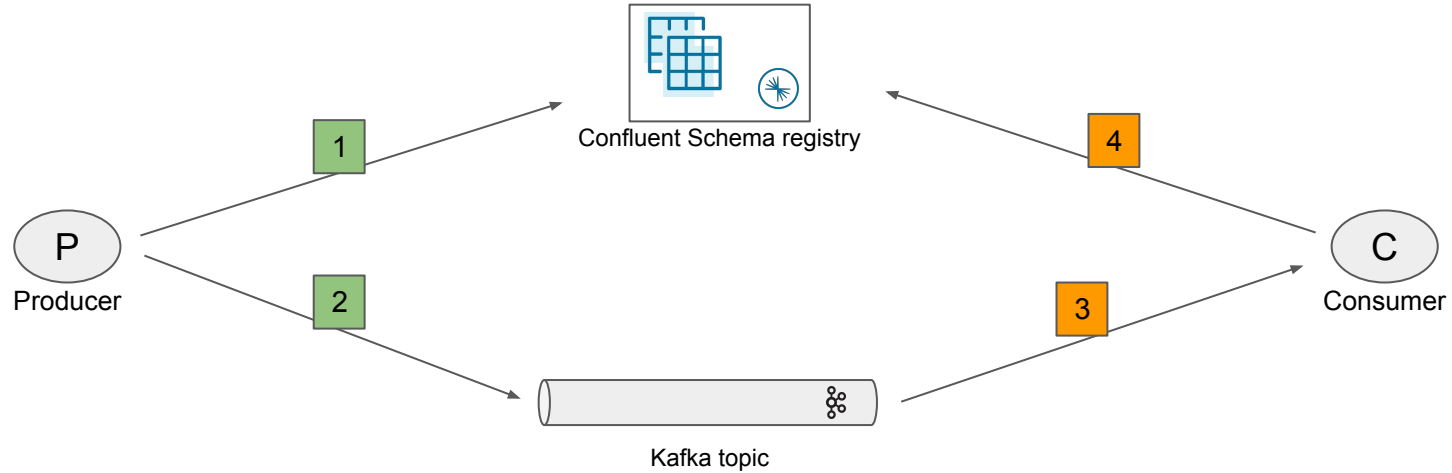
- Schema Registry
  - How does Schema registry works?
  - How YANG push modules are structured
  - Design choices for YANG integration
- YangKit
  - Interface with YangKit (JSON and CBOR)
  - Content encapsulated in “data” node
  - YANG data validation
  - YANG push Notification validation
- Missing pieces: libyangserde (C library adding the MAGIC BYTE and the schema Id into the Kafka message)

# Schema Registry

# Full Implementation-level architecture

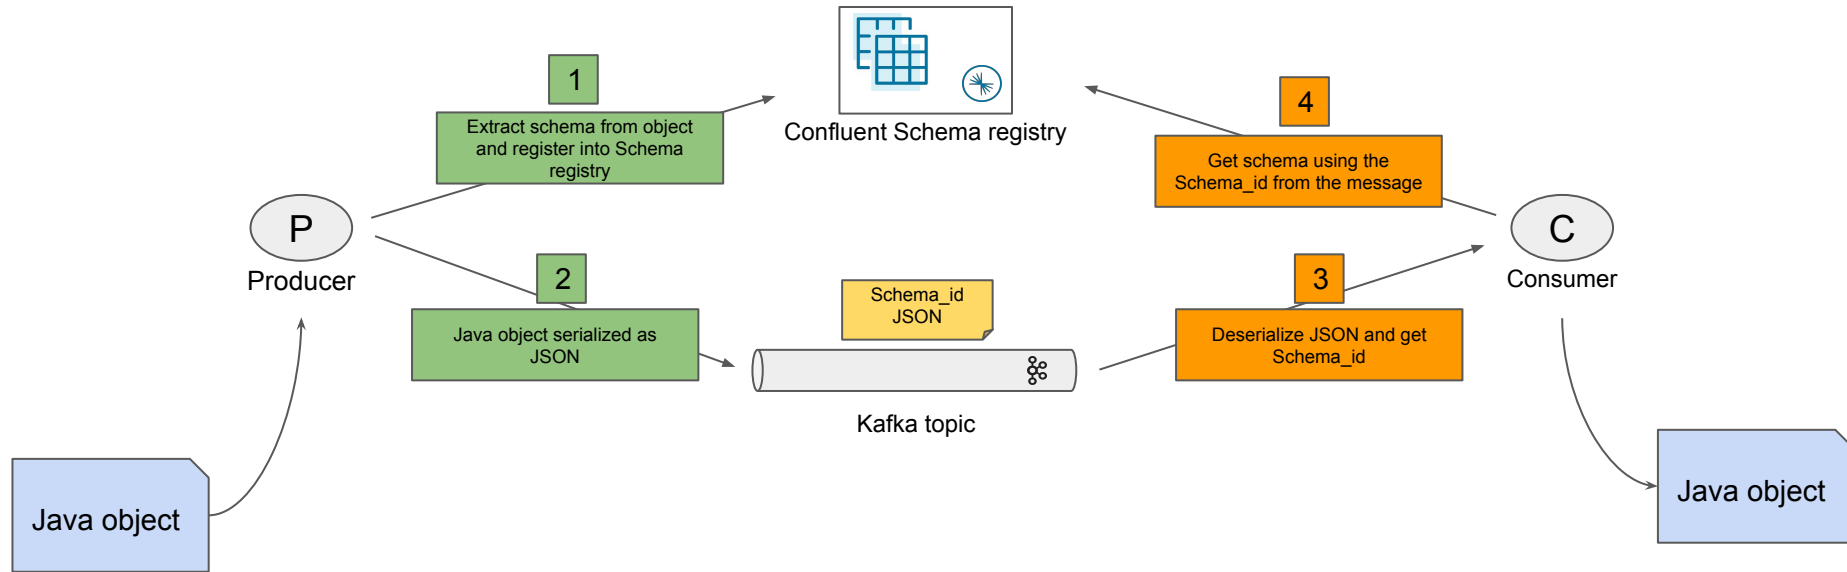


# How does Schema registry works?

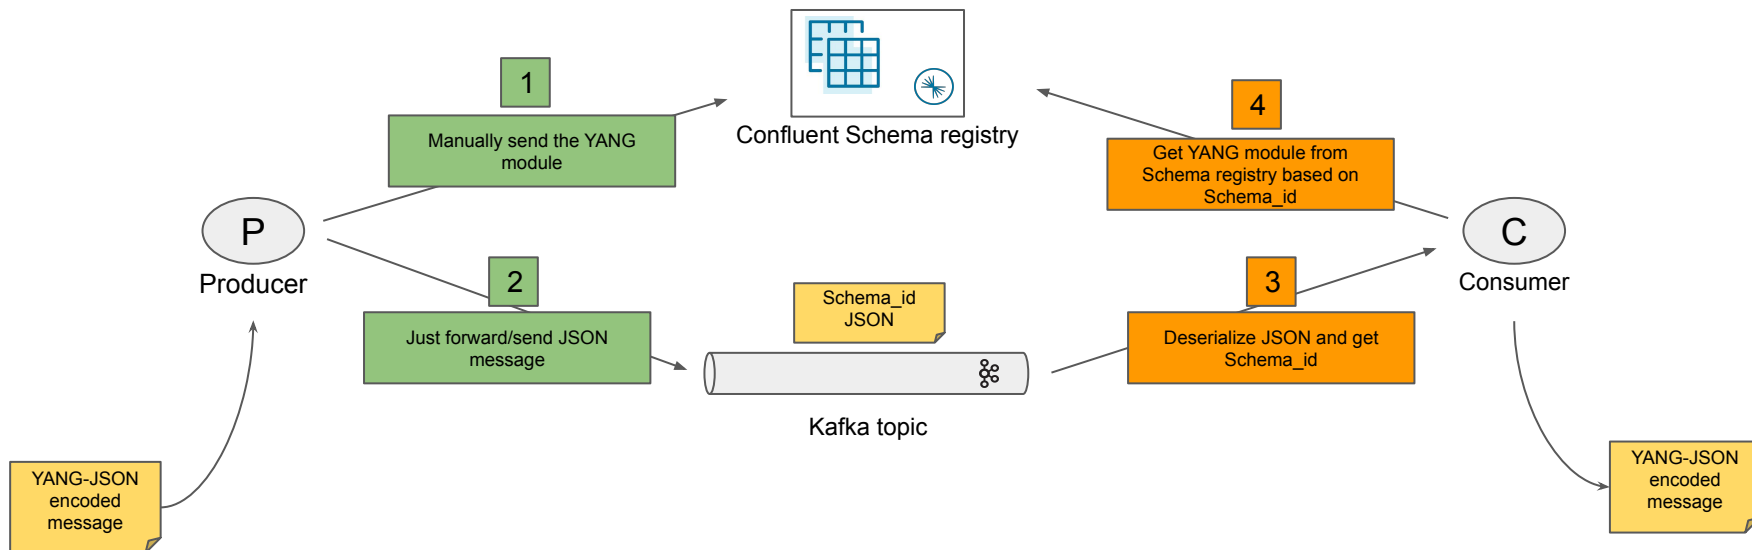


1. register schema
2. Send <MAGICBYTE, Schema ID, Content>
3. Receive <MAGICBYTE, Schema ID, Content>
4. Get schema

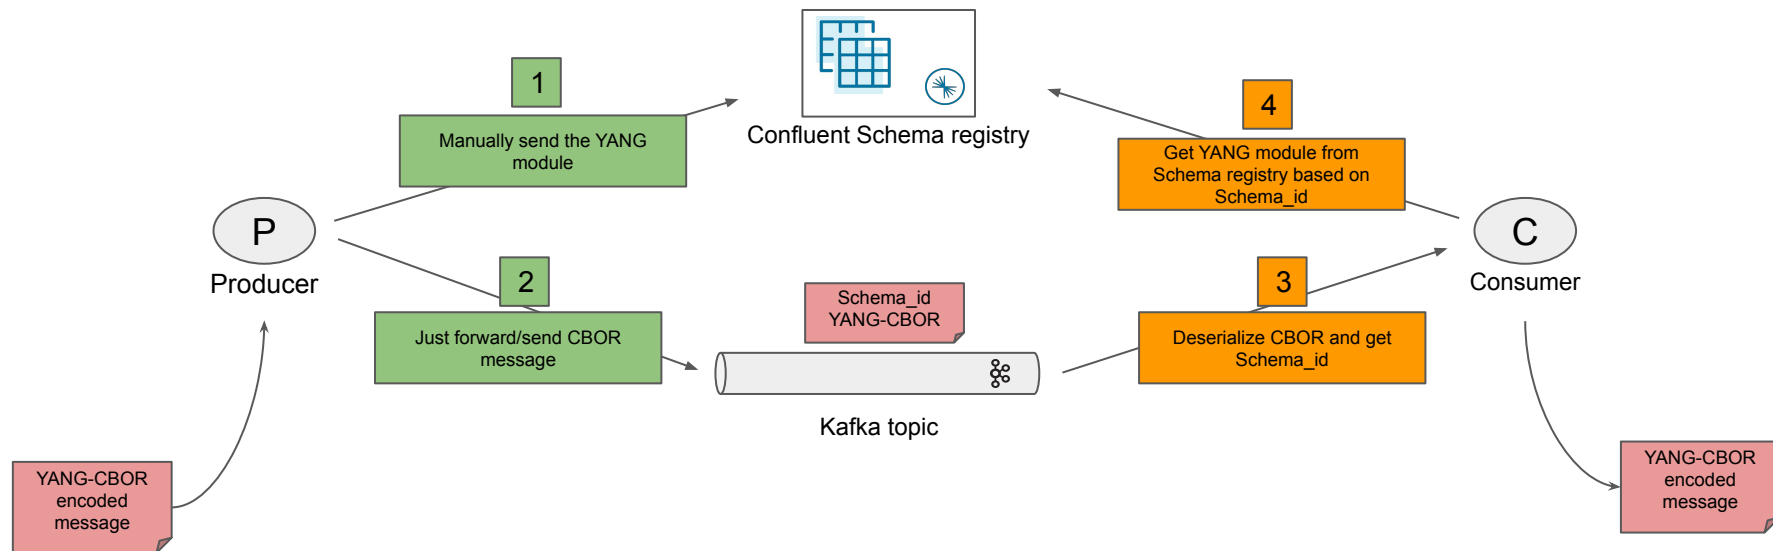
# Main usecase for Schema registry: Plain Old Java Objects



# Main usecase for YANG: Using YANG-JSON (without POJO)



# Main usecase for YANG: Using YANG-CBOR (without POJO)



*Same as YANG-JSON encoding but in YANG-CBOR*



# Nature of YANG modules / schemas

Example:

```
module insa-container {  
  yang-version 1.1;  
  namespace "urn:ietf:params:xml:ns:yang:insa-container";  
  prefix ic;  
  
  container first-container {  
    config false;  
  
    leaf address {  
      type string;  
    }  
    leaf port {  
      type uint8;  
    }  
  }  
}
```

```
module insa-augment {  
  yang-version 1.1;  
  namespace "urn:ietf:params:xml:ns:yang:insa-augment";  
  prefix ia;  
  
  import insa-container {  
    prefix ic;  
    reference "RFC XXXX: YYYY";  
  }  
  
  augment "/ic:first-container" {  
    leaf mtu {  
      type string;  
      config false;  
    }  
    container second-container {  
      config false;  
      leaf identifier {  
        type string;  
      }  
    }  
  }  
}
```

```
module insa-augment-bis {  
  yang-version 1.1;  
  namespace "urn:ietf:params:xml:ns:yang:insa-augment-bis";  
  prefix iab;  
  
  import insa-container {  
    prefix ic;  
    reference "RFC XXXX: YYYY";  
  }  
  import insa-augment {  
    prefix ia;  
    reference "RFC XXXX: YYYY";  
  }  
  
  augment "/ic:first-container/ia:second-container" {  
    leaf name {  
      type string;  
      config false;  
    }  
  }  
}
```

# Nature of YANG modules / schemas → Dependencies

```
module insa-container {  
  yang-version 1.1;  
  namespace "urn:ietf:params:xml:ns:yang:insa-container";  
  prefix ic;  
  
  container first-container {  
    config false;  
  
    leaf address {  
      type string;  
    }  
    leaf port {  
      type uint8;  
    }  
  }  
}
```

```
module insa-augment {  
  yang-version 1.1;  
  namespace "urn:ietf:params:xml:ns:yang:insa-augment";  
  prefix ia;  
  
  import insa-container {  
    prefix ic;  
    reference "RFC XXXX: YYYY";  
  }  
  
  augment "/ic:first-container" {  
    leaf mtu {  
      type string;  
      config false;  
    }  
    container second-container {  
      config false;  
      leaf identifier {  
        type string;  
      }  
    }  
  }  
}
```

```
module insa-augment-bis {  
  yang-version 1.1;  
  namespace "urn:ietf:params:xml:ns:yang:insa-augment-bis";  
  prefix iab;  
  
  import insa-container {  
    prefix ic;  
    reference "RFC XXXX: YYYY";  
  }  
  import insa-augment {  
    prefix ia;  
    reference "RFC XXXX: YYYY";  
  }  
  
  augment "/ic:first-container/ia:second-container" {  
    leaf name {  
      type string;  
      config false;  
    }  
  }  
}
```

insa-container.yang ← insa-augment.yang ← insa-augment-bis.yang

# Nature of YANG modules / schemas: registration in Schema registry

insa-container.yang ← insa-augment.yang ← insa-augment-bis.yang

1

2

3

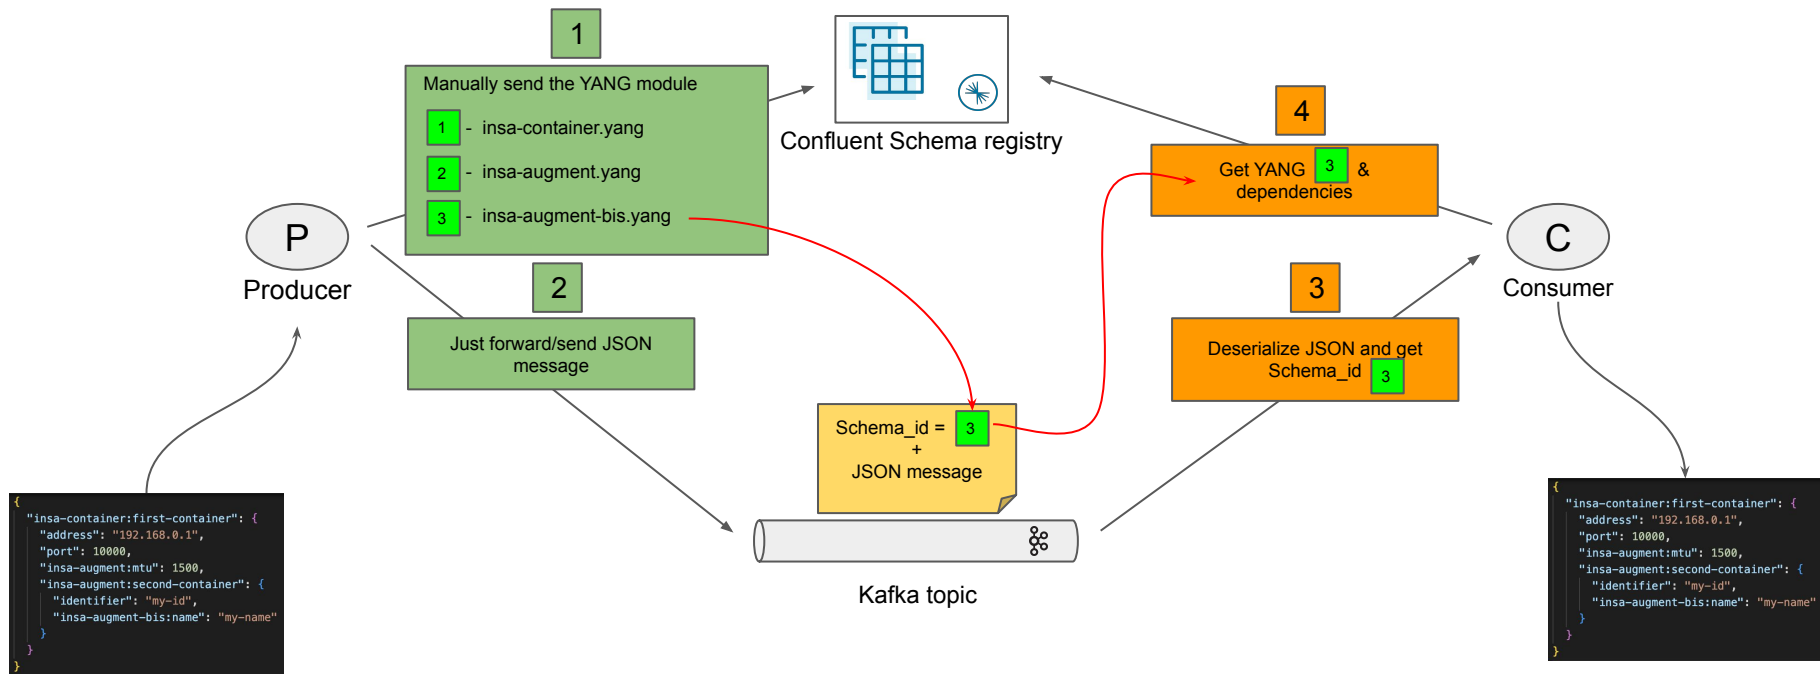
1. Register insa-container.yang (subject **insa-container**)
2. Register insa-augment.yang (subject **insa-augment**)
3. Register insa-augment-bis.yang (subject **insa-augment-bis**)
4. Use last subject (*insa-augment-bis*) to validate message

```
module: insa-container
  +-ro first-container
    +-ro address?      string
    +-ro port?         uint8
    +-ro ia:mtu?       string
    +-ro ia:second-container
      +-ro ia:identifier? string
      +-ro iab:name?    string
```



```
{
  "insa-container:first-container": {
    "address": "192.168.0.1",
    "port": 10000,
    "insa-augment:mtu": 1500,
    "insa-augment:second-container": {
      "identifier": "my-id",
      "insa-augment-bis:name": "my-name"
    }
  }
}
```

# Main usecase for YANG: Using YANG-JSON (without POJO)



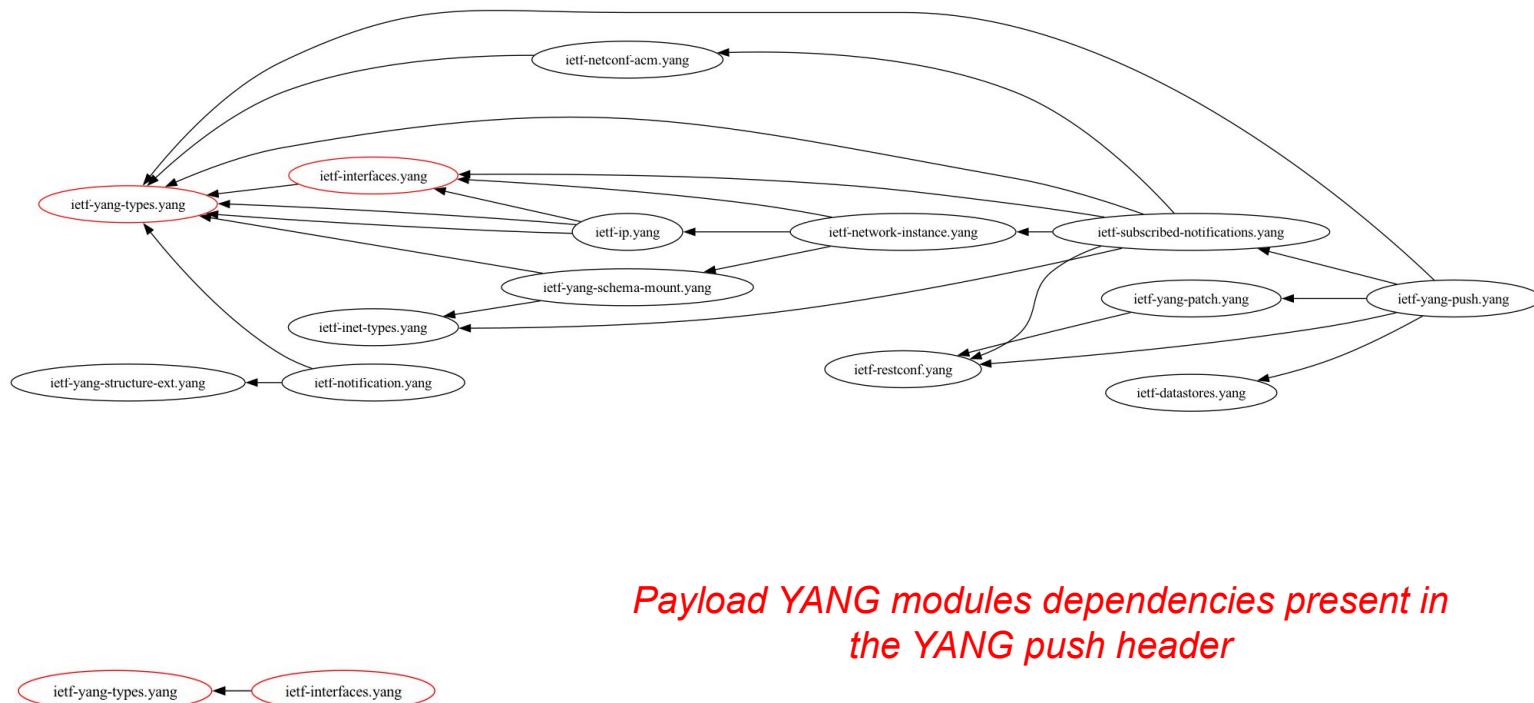
# YANG push in Schema registry: **example 1**

```
{
  "ietf-notification:notification": {
    "eventTime": "2023-03-25T08:30:11.22Z",
    "ietf-notification-sequencing:sysName": "example-router",
    "ietf-notification-sequencing:sequenceNumber": 1,
    "ietf-yang-push:push-update": {
      "id": 6666,
      "ietf-yang-push-netobs-timestamping:observation-time": "2023-03-25T08:30:11.22Z",
      "datastore-contents": {
        "ietf-interfaces:interfaces": [
          {
            "interface": {
              "name": "eth0",
              "type": "iana-if-type:ethernetCsmacd",
              "oper-status": "up"
            }
          }
        ]
      }
    }
  }
}
```

YANG-push header

YANG-push payload

## YANG push in Schema registry: dependencies



## YANG push in Schema registry: **example 2**

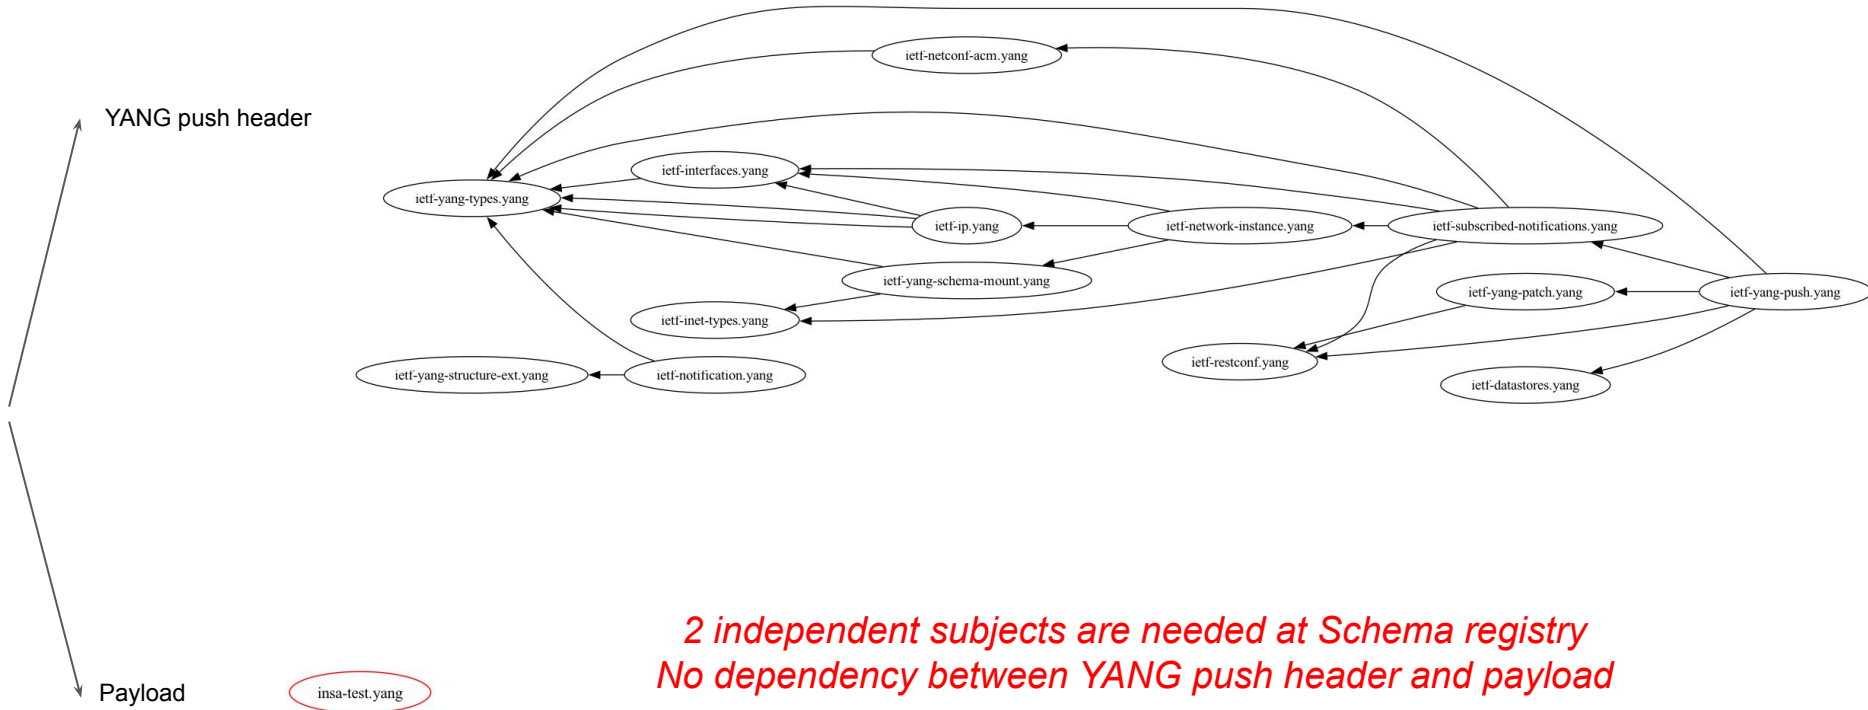
```
{
  "ietf-notification:notification": {
    "eventTime": "2023-03-25T08:30:11.22Z",
    "ietf-notification-sequencing:sysName": "example-router",
    "ietf-notification-sequencing:sequenceNumber": 1,
    "ietf-yang-push:push-update": {
      "id": 6666,
      "ietf-yang-push-netobs-timestamping:observation-time": "2023-03-25T08:30:11.22Z",
      "datastore-contents": {
        "insa-test:insa-container": {
          "computer": "my-computer",
          "router": 26
        }
      }
    }
  }
}
```

YANG-push header

Independent YANG

**Note:** *insa-test.yang* is a non-standard YANG module for this example. Same usecase as if OpenConfig YANG modules were used.

# YANG push in Schema registry: dependencies





# Designs for integrating YANG to Schema registry

- (1) Create a YANG module integrating all the YANG push header and payload
  - First approach proposed at Hackathon 117
  - **Issue**: Creating a new YANG module changes the namespace of the encoded message.
  
- (2) Register all the YANG modules to a Schema Context
  - Approach taken by YangKit (currently testing this approach at INSA)
  - **Issue**: Schema registry needs to support “one subject is associated to multiple models” (not only augmentations, but independent modules)

# Impact on Schema registry using Design (2)

## Register all the YANG modules to a Schema Context

- **Issue:** Schema registry needs to support “one subject is associated to multiple models” (not only augmentations, but independent modules)
  - Solution 1: Implement support in schema registry
    - API is changed and not “compliant” to Confluent approach (meaning we are implementing a new API)
  - Solution 2: (workaround) Global SchemaContext per Schema Registry
    - API is not changed, but upon request, the global SchemaContext is used
    - YANG versioning and BC/NBC checks cannot be supported

# Yangkit

Interfaces

Gaps

# Interface with YangKit (JSON)

Input:

- JsonNode (Jackson library)
- SchemaContext (Class having all Serialised yang modules)

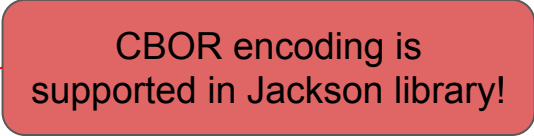
Output:

- Notification is valid/invalid
- Data is valid/invalid

# Interface with YangKit (CBOR)

Input:

- JsonNode (Jackson library)
- SchemaContext (Class having all Serialised yang modules)



CBOR encoding is supported in Jackson library!

Output:

- Notification is valid/invalid
- Data is valid/invalid

# Validating YANG-JSON **data**

YANG data is wrapped in “data” node

- *Is it a Standard/expected behavior?*
- *Behavior coming from NETCONF RFC6241?*
- Easy to change

```
{  
  "data": {  
    "insa-custom:insa-container": {  
      "computer": 1,  
      "router": 234  
    }  
  }  
}
```

# Yangkit validation gaps

- Type validation
- Mandatory
- Lists
- Unknown elements

# Validating YANG-JSON data: Type validation

Leaf **computer**:

- YANG definition: string
- Content in the data: integer

→ Validator doesn't throw any error


```
{
  "data": {
    "insa-custom:insa-container": {
      "computer": 1,
      "router": 234
    }
  }
}
```



```
description
  "insa-test YANG module.";

revision 2023-09-05 {
  description "Initial version.";
}

container insa-container {
  config false;
  leaf computer {
    type string;
    mandatory true;
    description "computer";
  }
  leaf router {
    type uint8;
    description "router";
  }
}
```





# Validating YANG-JSON data: Mandatory leaves


Leaf **computer**:

- YANG definition: **mandatory**
- Content in the data: **not present**

→ Validator doesn't throw any error



```
{
  "data": {
    "insa-custom:insa-container": {
      "router": 234
    }
  }
}
```



```
description
  "insa-test YANG module.";

revision 2023-09-05 {
  description "Initial version.";
}

container insa-container {
  config false;
  leaf computer {
    type string;
    mandatory true;
    description "computer";
  }
  leaf router {
    type uint8;
    description "router";
  }
}
```

# Validating YANG-JSON data: Lists (missing keys)

Leaf **computer**:


- YANG definition: **Key of a list**
- Content in the data: **not present**

→ Validator throws a “**missing-element**” error

```
{
  "data": {
    "insa-custom:insa-container": [{
      "router": 234
    }]
  }
}
```



```
list insa-container {
  key computer;
  config false;
  leaf computer {
    type string;
    mandatory true;
    description "computer";
  }
  leaf router {
    type uint8;
    description "router";
  }
}
```



# Validating YANG-JSON data: Unknown elements

Leaf `invalid_key`:

- YANG definition: **Not defined**
- Content in the data: **Present**

→ Validator throws an “**unknown element**” error



```
{
  "data": {
    "insa-custom:insa-container": {
      "computer": "computer",
      "router": 234,
      "invalid_key": "hey"
    }
  }
}
```

```
description
  "insa-test YANG module.";

revision 2023-09-05 {
  description "Initial version.";
}

container insa-container {
  config false;
  leaf computer {
    type string;
    mandatory true;
    description "computer";
  }
  leaf router {
    type uint8;
    description "router";
  }
}
```

# Validating YANG push Notification

- Accepts a YANG module defining a structure
- Can separate YANG push header from the payload
- Data Validation: similar to YANG data

# Corner cases to have in mind

## XPath is not managed by Yangkit

- YANG push subscription need to validate the data defined by the XPath
- The rest of the YANG module, the data should not be present

→ Yangkit approach: validate against a general YANG tree defined by **SchemaContext**

# Corner cases to have in mind

Xpath=/example:insa-container

Xpath=/example:test-container

```
revision 2019-09-09 {
  description "Initial version.";
}

container insa-container {
  config false;
  leaf computer {
    type string;
    description "computer";
  }
  leaf router {
    type uint8;
    description "router";
  }
}

container test-container {
  config false;
  leaf my-id {
    type string;
    mandatory true;
    description "identifier";
  }
  leaf subscription {
    type string;
    description "subscription";
  }
}
```

Mandatory field

# Corner cases to have in mind

Subscription to

- Xpath=/example:insa-container

Yangkit default behavior (Checked with Frank):

- Data not having /test-container/my-id should **fail**

```
{
  "data": {
    "insa-custom:insa-container": {
      "computer": "computer",
      "router": 234
    }
  }
}
```

```
revision 2019-09-09 {
  description "Initial version.";
}
```

```
container insa-container {
  config false;
  leaf computer {
    type string;
    description "computer";
  }
  leaf router {
    type uint8;
    description "router";
  }
}

container test-container {
  config false;
  leaf my-id {
    type string;
    mandatory true;
    description "identifier";
  }
  leaf subscription {
    type string;
    description "subscription";
  }
}
```

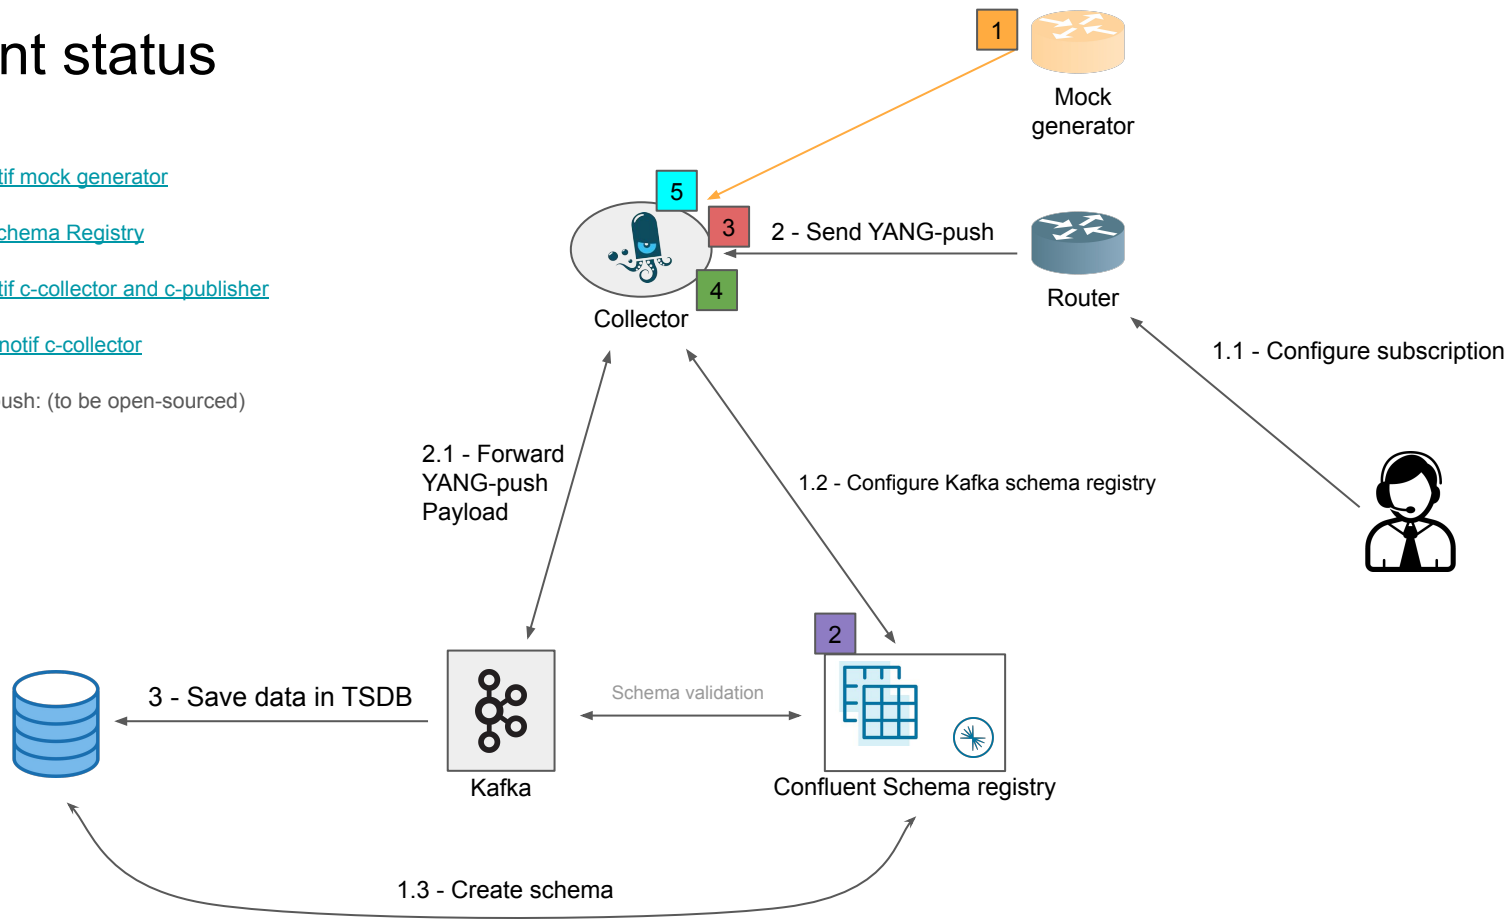
# Missing pieces

libyangserdes



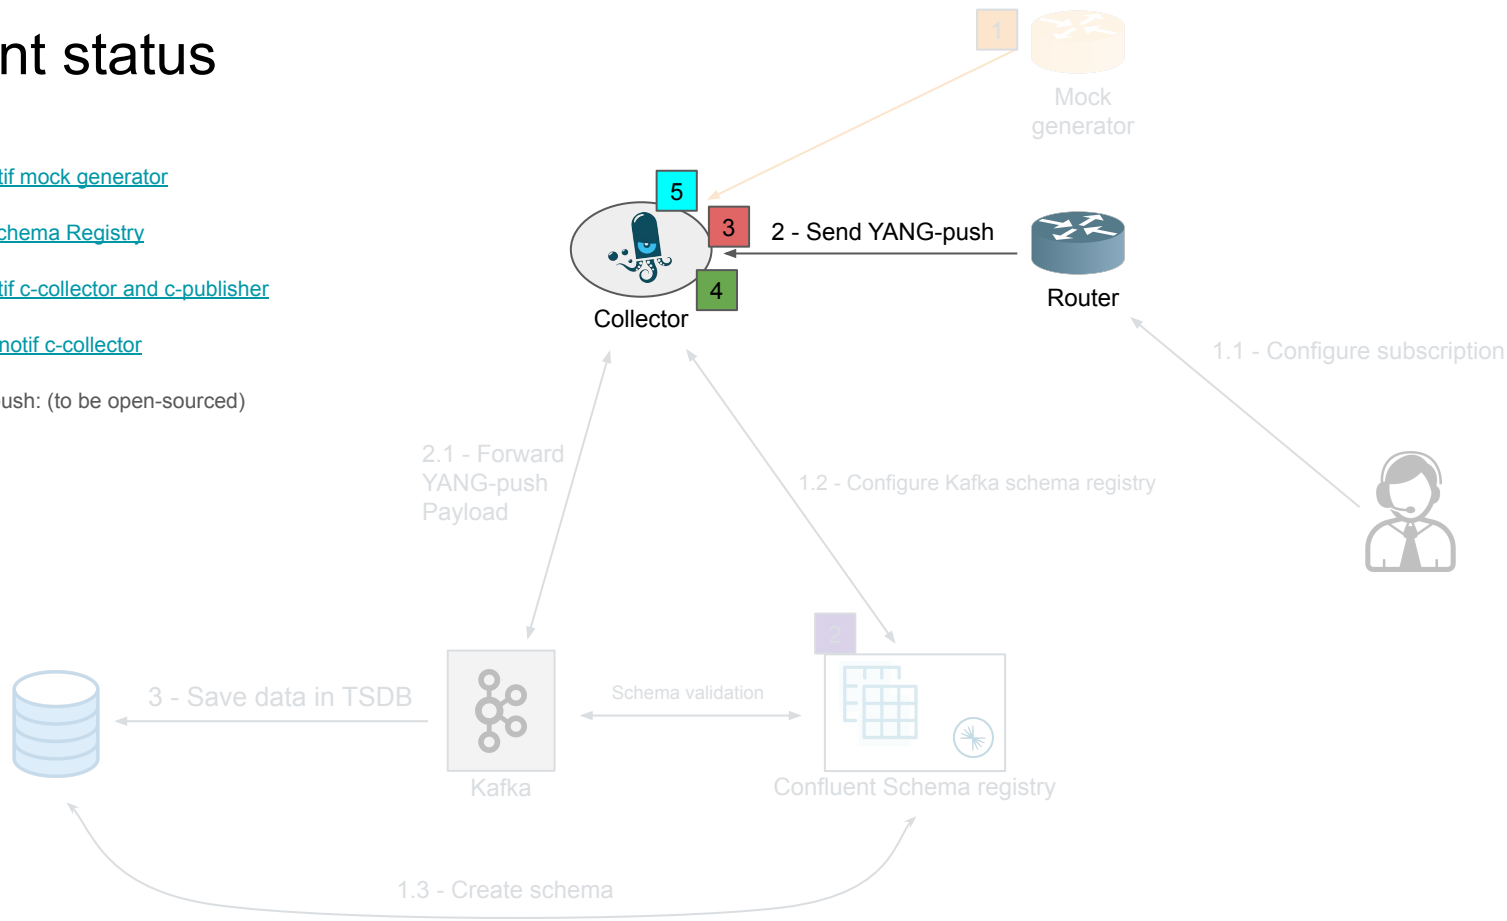
# Current status

- 1 [UDP-notif mock generator](#)
- 2 [Kafka Schema Registry](#)
- 3 [UDP-notif c-collector and c-publisher](#)
- 4 [HTTPS-notif c-collector](#)
- 5 libyangpush: (to be open-sourced)

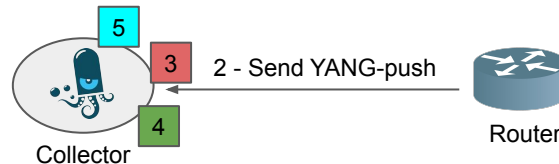


# Current status

- 1 [UDP-notif mock generator](#)
- 2 [Kafka Schema Registry](#)
- 3 [UDP-notif c-collector and c-publisher](#)
- 4 [HTTPS-notif c-collector](#)
- 5 libyangpush: (to be open-sourced)



## Current status: router - collector



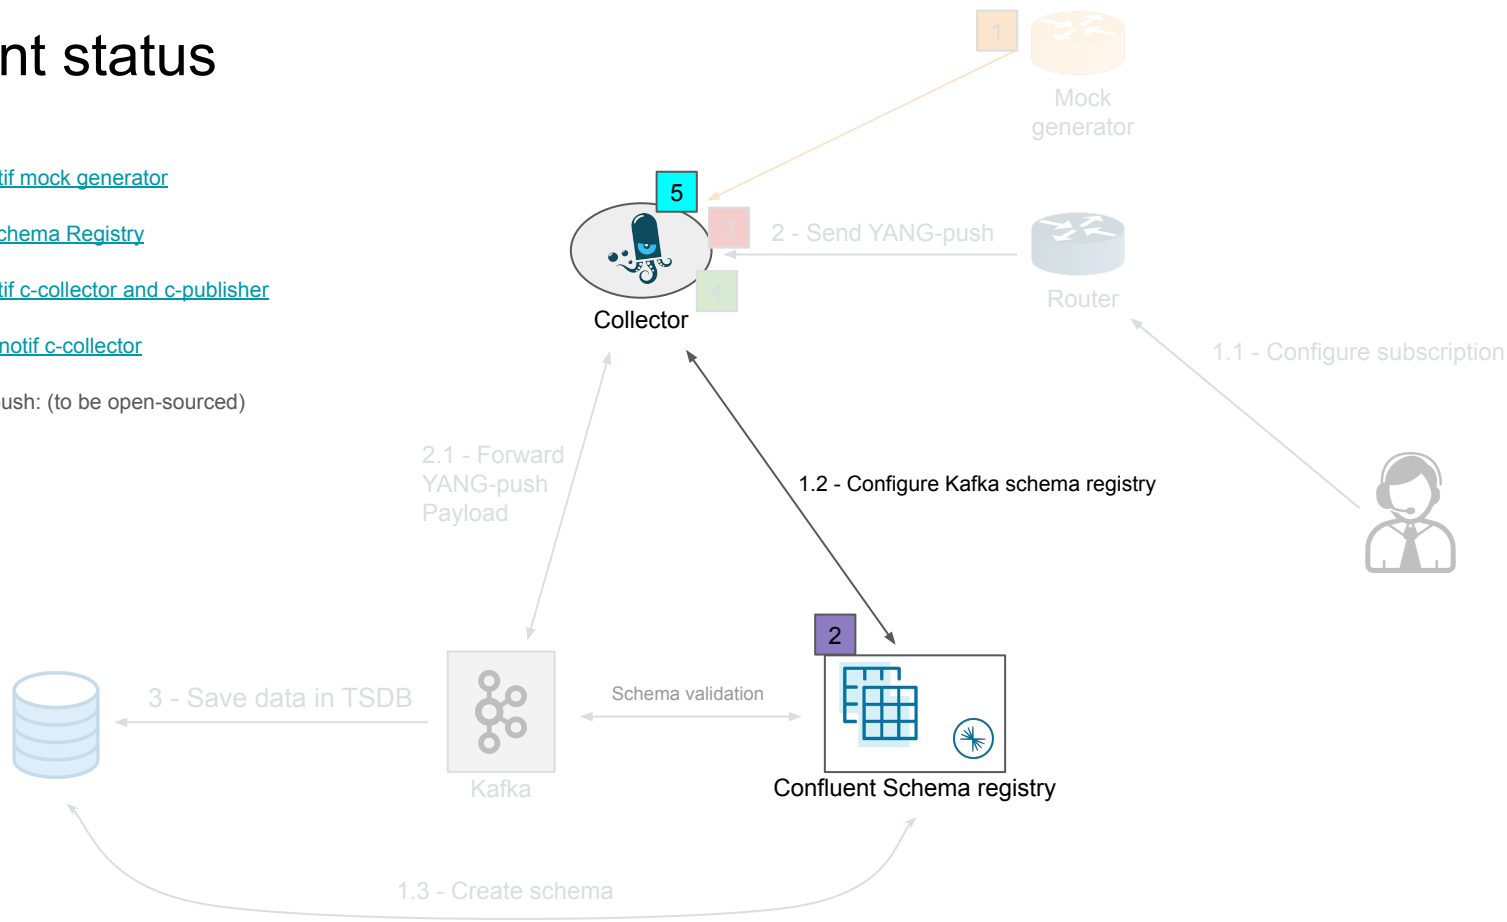
- Router
  - able to craft UDP-notif messages using: [UDP-notif c-collector and c-publisher](#)
- Collector
  - able to collect UDP-notif messages using: [UDP-notif c-collector and c-publisher](#)
  - able to collect HTTPS-notif messages using: [HTTPS-notif c-collector](#)
  - able to get YANG modules dependencies using: libyangpush (to be open-sourced)

*YANG-JSON and YANG-CBOR encoding is generated at Router independently from the transport protocol.*

*Current Transports already support **YANG-JSON** and **YANG-CBOR** MediaType*

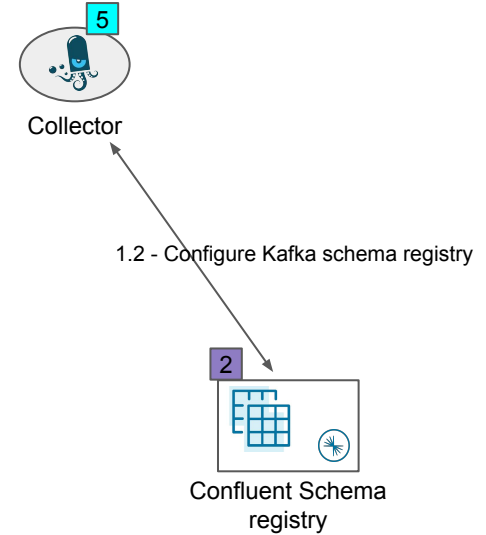
# Current status

- 1 [UDP-notif mock generator](#)
- 2 [Kafka Schema Registry](#)
- 3 [UDP-notif c-collector and c-publisher](#)
- 4 [HTTPS-notif c-collector](#)
- 5 libyangpush: (to be open-sourced)



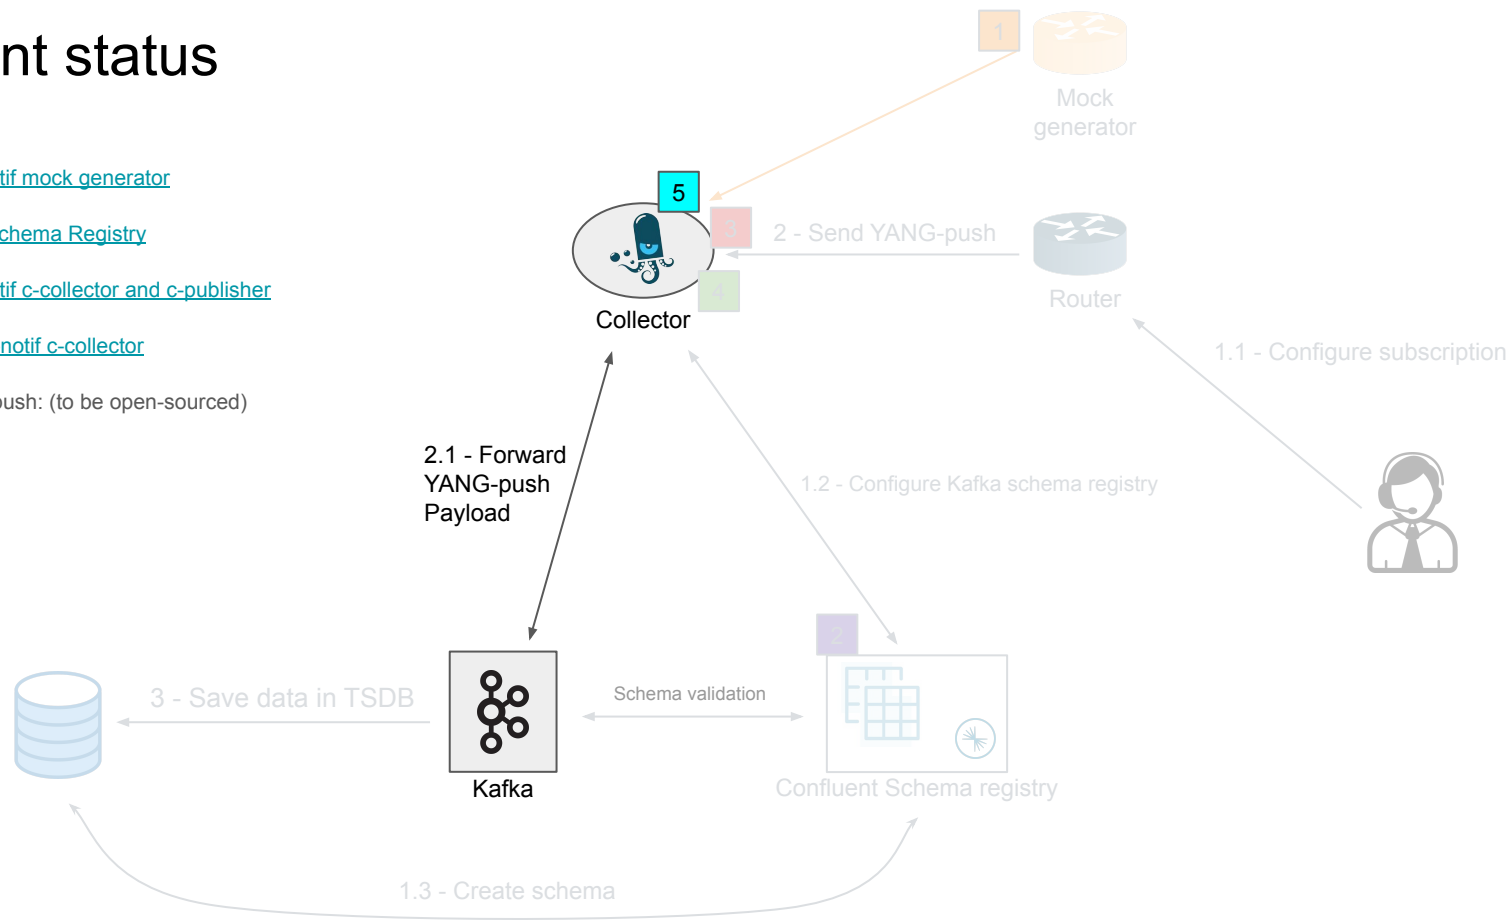
# Current status: collector - schema registry

- Collector
  - when the YANG push message is received, we are able to get all YANG modules from the router using `libyangpush`
  - Generate registration to Schema Registry using `libyangpush`
- Schema registry
  - Support **YANG** as a schema type
  - *WIP: one subject defines one YANG module or one Schema Context*



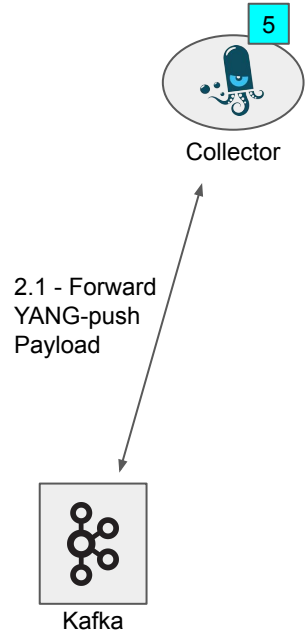
# Current status

- 1 [UDP-notif mock generator](#)
- 2 [Kafka Schema Registry](#)
- 3 [UDP-notif c-collector and c-publisher](#)
- 4 [HTTPS-notif c-collector](#)
- 5 libyangpush: (to be open-sourced)



# Current status: collector - Kafka topic

- Collector
  - Once the collector got the Schema\_id serialize the message to Kafka using **libyangserdes** (Craft Kafka message with MAGICBYTE and Schema\_id)
- Kafka
  - is able to receive JSON or CBOR messages



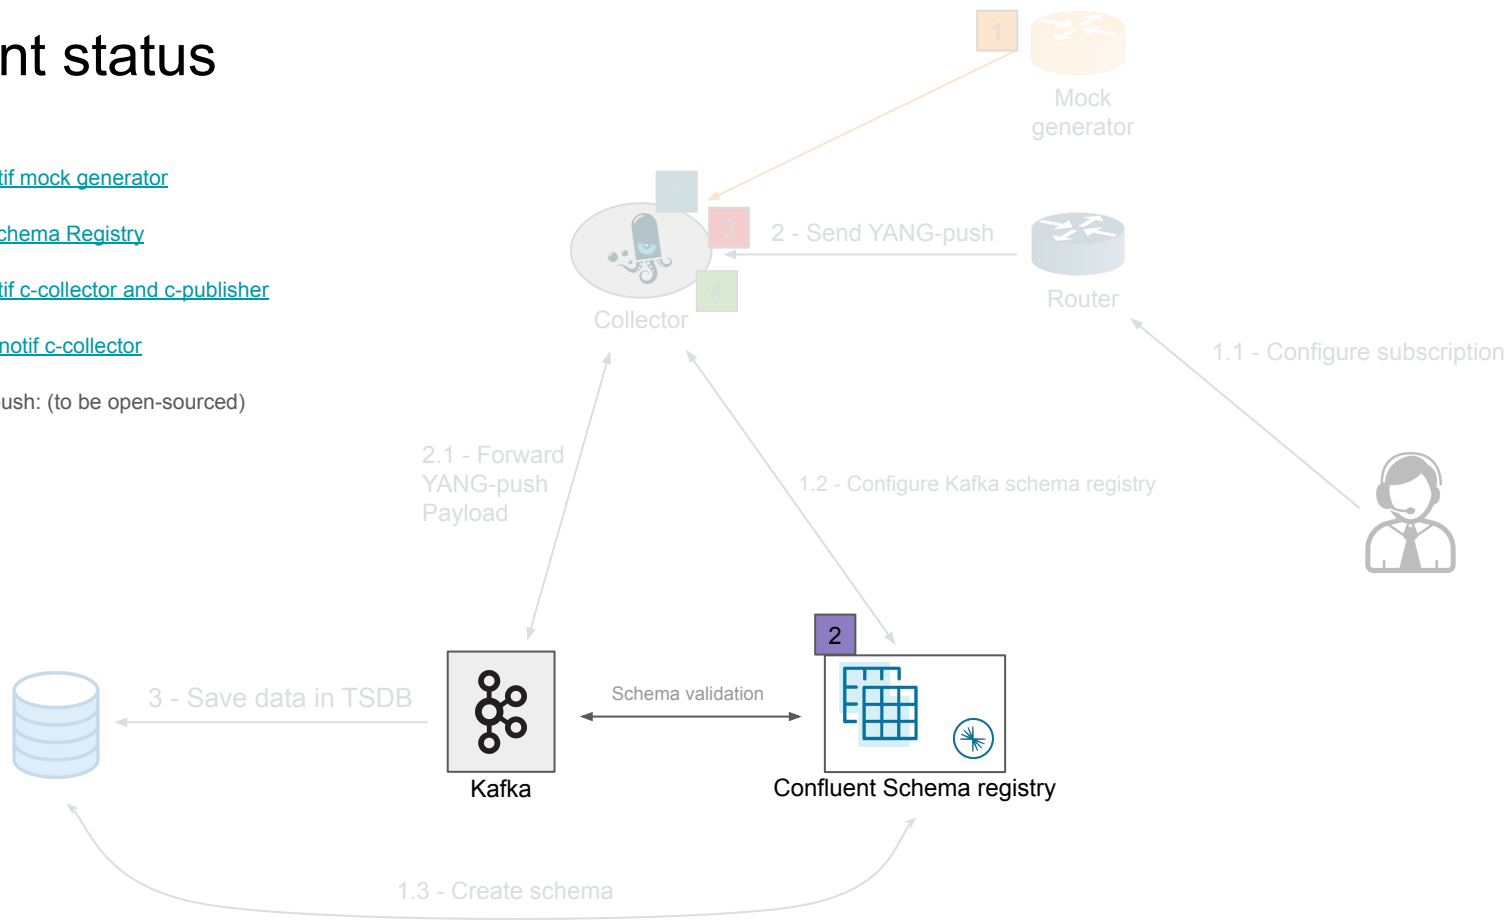
# libyangserdes

- Client uses `libyangpush` to register all the YANG modules to Schema Registry
- Client uses `libyangserdes` to craft the kafka serialised message including the ***MAGICBYTE*** and the ***Schema\_id***
- `libyangserdes` must not modify the content of the message

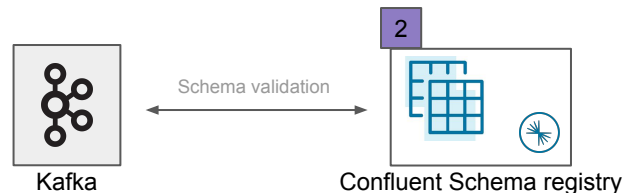


# Current status

- 1 [UDP-notif mock generator](#)
- 2 [Kafka Schema Registry](#)
- 3 [UDP-notif c-collector and c-publisher](#)
- 4 [HTTPS-notif c-collector](#)
- 5 libyangpush: (to be open-sourced)



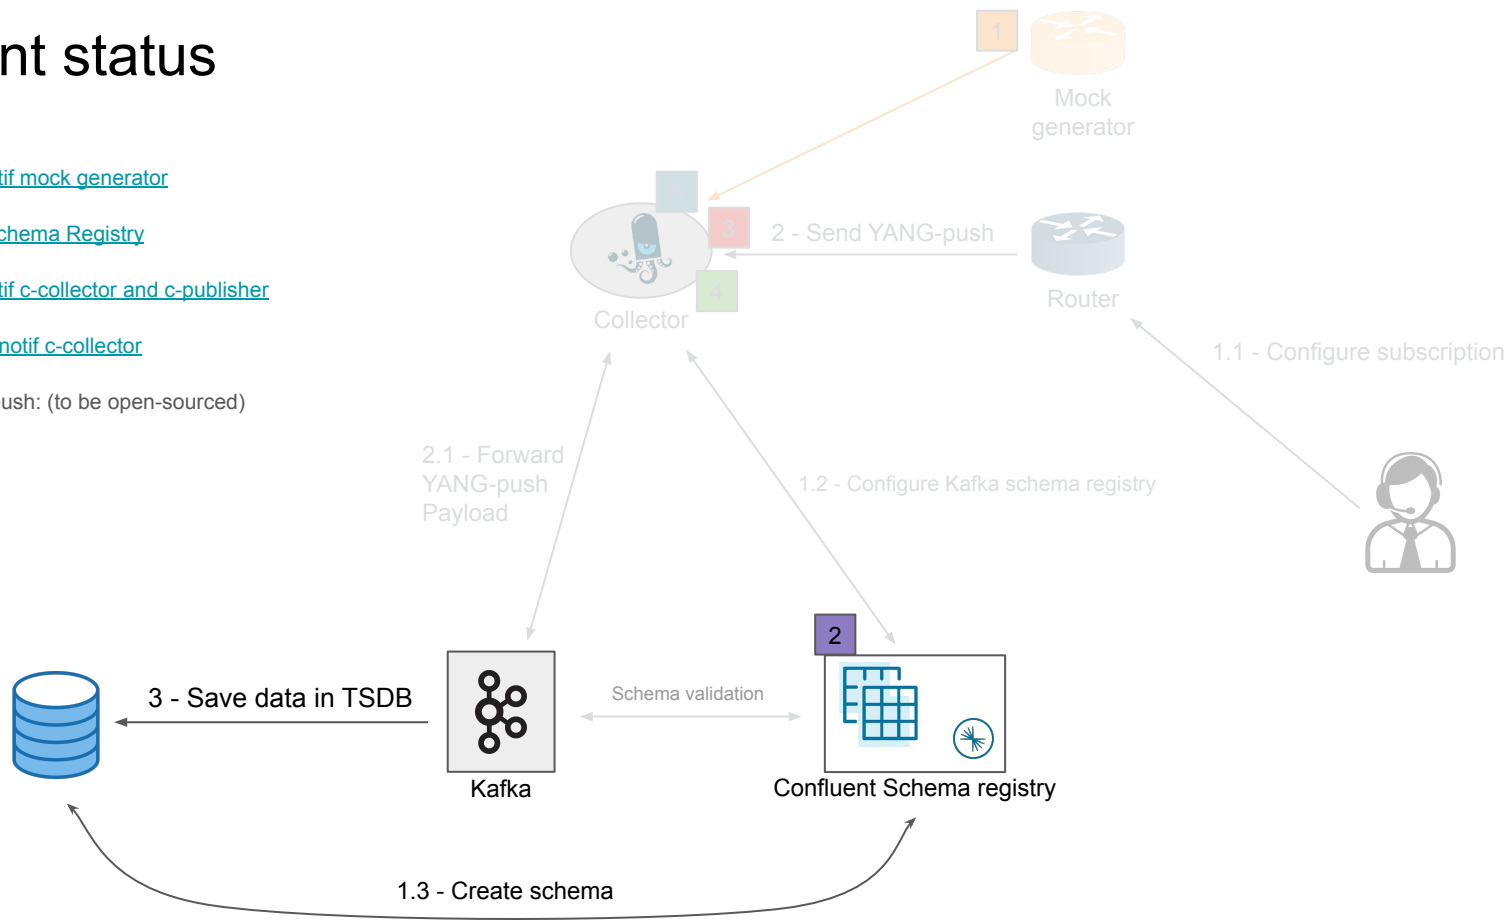
# Current status: Kafka - Schema Registry



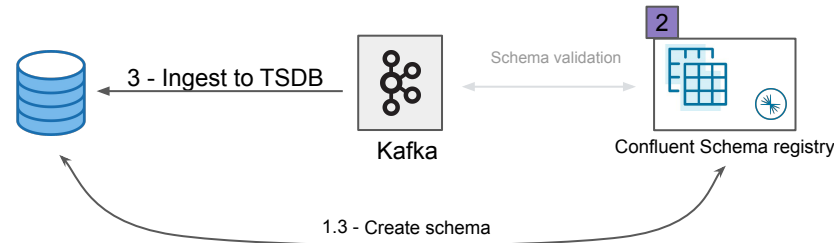
- Kafka
  - is able to get the Schema\_id from the message
  - is able to request the schema from the Schema Registry and validate the content
  - **WIP: YANG validation with Yangkit**
- Confluent Schema Registry
  - is able to provide the schema / schema context
  - **WIP: Current discussions if current API need to be modified to accommodate YANG subjects**

# Current status

- 1 [UDP-notif mock generator](#)
- 2 [Kafka Schema Registry](#)
- 3 [UDP-notif c-collector and c-publisher](#)
- 4 [HTTPS-notif c-collector](#)
- 5 libyangpush: (to be open-sourced)



# Current status: TSDB ingestion



- TSDB
  - Druid uses Kafka connect to create/ingest data to the database
  - Missing: YANG connector (part of Druid)
- Kafka
  - is able to provide the schema / schema context
  - WIP: Current discussions if current API need to be modified to accommodate YANG subjects
- Confluent Schema registry
  - is able to provide the schema / schema context
  - WIP: Current discussions if current API need to be modified to accommodate YANG subjects